

Loan Approval Prediction Model

Problem Statement:

Develop a machine learning model to predict **loan approval status** based on applicant features. The model will be trained on a dataset containing loan applications and their outcomes. It will help financial institutions automate and improve the accuracy of loan approval decisions.

Objective:

To predict **Loan Approval Status** using a Machine Learning Model.

Benefits:

The benefits of this solution include:

- **Lenders** can make **faster and more accurate loan approval decisions**, reducing manual effort.
- **Applicants** will get **quicker responses** and **fairer evaluations** based on objective data.
- **Financial institutions** can **minimize default risks** by identifying high-risk applicants more effectively.

#Importing Required Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix, roc_curve, roc_auc_score
```

Load the Dataset

```
df = pd.read_csv('loan_approval_dataset.csv')
df.head()
```

	loan_id	no_of_dependents	education	self_employed	
0	9600000	1	2	Graduate	No
1	4100000	2	0	Not Graduate	Yes
2	9100000	3	3	Graduate	No

3	4	3	Graduate	No
8200000				
4	5	5	Not Graduate	Yes
9800000				

	loan_amount	loan_term	cibil_score	
residential_assets_value \				
0	29900000	12	778	2400000
1	12200000	8	417	2700000
2	29700000	20	506	7100000
3	30700000	8	467	18200000
4	24200000	20	382	12400000

	commercial_assets_value	luxury_assets_value	
bank_asset_value \			
0	17600000	22700000	8000000
1	2200000	8800000	3300000
2	4500000	33300000	12800000
3	3300000	23300000	7900000
4	8200000	29400000	5000000

	loan_status
0	Approved
1	Rejected
2	Rejected
3	Rejected
4	Rejected

Getting shape of dataset

df.shape

(4269, 13)

df.head()

	loan_id	no_of_dependents	education	self_employed
income_annum \				
0	1	2	Graduate	No
9600000				
1	2	0	Not Graduate	Yes
4100000				

2	3	3	Graduate	No
9100000				
3	4	3	Graduate	No
8200000				
4	5	5	Not Graduate	Yes
9800000				

	loan_amount	loan_term	cibil_score	
	residential_assets_value \			
0	29900000	12	778	2400000
1	12200000	8	417	2700000
2	29700000	20	506	7100000
3	30700000	8	467	18200000
4	24200000	20	382	12400000

	commercial_assets_value	luxury_assets_value	
	bank_asset_value \		
0	17600000	22700000	8000000
1	2200000	8800000	3300000
2	4500000	33300000	12800000
3	3300000	23300000	7900000
4	8200000	29400000	5000000

	loan_status
0	Approved
1	Rejected
2	Rejected
3	Rejected
4	Rejected

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 4269 entries, 0 to 4268
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	loan_id	4269 non-null	int64
1	no_of_dependents	4269 non-null	int64
2	education	4269 non-null	object
3	self_employed	4269 non-null	object

4	income_annum	4269	non-null	int64
5	loan_amount	4269	non-null	int64
6	loan_term	4269	non-null	int64
7	cibil_score	4269	non-null	int64
8	residential_assets_value	4269	non-null	int64
9	commercial_assets_value	4269	non-null	int64
10	luxury_assets_value	4269	non-null	int64
11	bank_asset_value	4269	non-null	int64
12	loan_status	4269	non-null	object

dtypes: int64(10), object(3)

memory usage: 433.7+ KB

None

Statistical summary of numerical columns

df.describe()

	loan_id	no_of_dependents	income_annum	loan_amount	\
count	4269.000000	4269.000000	4.269000e+03	4.269000e+03	
mean	2135.000000	2.498712	5.059124e+06	1.513345e+07	
std	1232.498479	1.695910	2.806840e+06	9.043363e+06	
min	1.000000	0.000000	2.000000e+05	3.000000e+05	
25%	1068.000000	1.000000	2.700000e+06	7.700000e+06	
50%	2135.000000	3.000000	5.100000e+06	1.450000e+07	
75%	3202.000000	4.000000	7.500000e+06	2.150000e+07	
max	4269.000000	5.000000	9.900000e+06	3.950000e+07	

	loan_term	cibil_score	residential_assets_value	\
count	4269.000000	4269.000000	4.269000e+03	
mean	10.900445	599.936051	7.472617e+06	
std	5.709187	172.430401	6.503637e+06	
min	2.000000	300.000000	-1.000000e+05	
25%	6.000000	453.000000	2.200000e+06	
50%	10.000000	600.000000	5.600000e+06	
75%	16.000000	748.000000	1.130000e+07	
max	20.000000	900.000000	2.910000e+07	

	commercial_assets_value	luxury_assets_value
bank_asset_value		
count	4.269000e+03	4.269000e+03
4.269000e+03		
mean	4.973155e+06	1.512631e+07
4.976692e+06		
std	4.388966e+06	9.103754e+06
3.250185e+06		
min	0.000000e+00	3.000000e+05
0.000000e+00		
25%	1.300000e+06	7.500000e+06
2.300000e+06		
50%	3.700000e+06	1.460000e+07
4.600000e+06		

```

75%          7.600000e+06          2.170000e+07
7.100000e+06
max          1.940000e+07          3.920000e+07
1.470000e+07

```

```
# Checking for Duplicate values
```

```
df.duplicated().sum()
```

```
np.int64(0)
```

```
# Dropping rows with null values
```

```
df.dropna(inplace=True)
```

```
#Removing Blank spaces from column names and values
```

```
df.columns = df.columns.str.strip()
```

```
df['loan_status'] = df['loan_status'].str.strip() # Remove spaces
```

```
df = df.apply(lambda x: x.str.strip() if x.dtype == 'object' else x)
```

```
# Data Transformation on Education, Self Employed Columns
```

```
df['education'] = df['education'].replace({'Not Graduate': 0,  
'Graduate': 1})
```

```
df['self_employed'] = df['self_employed'].replace({'No': 0, 'Yes': 1})
```

```
df.head()
```

	loan_id	no_of_dependents	education	self_employed	income_annum
0	1	2	1	0	9600000
1	2	0	0	1	4100000
2	3	3	1	0	9100000
3	4	3	1	0	8200000
4	5	5	0	1	9800000

	loan_amount	loan_term	cibil_score	residential_assets_value
0	29900000	12	778	2400000
1	12200000	8	417	2700000
2	29700000	20	506	7100000
3	30700000	8	467	18200000
4	24200000	20	382	12400000

	commercial_assets_value	luxury_assets_value	bank_asset_value
loan_status			

0	17600000	22700000	8000000
Approved			
1	2200000	8800000	3300000
Rejected			
2	4500000	33300000	12800000
Rejected			
3	3300000	23300000	7900000
Rejected			
4	8200000	29400000	5000000
Rejected			

Data Transformation on Loan Status Columns

```
df['loan_status'] = df['loan_status'].map({'Approved': 1, 'Rejected': 0})
```

```
df.head()
```

	loan_id	no_of_dependents	education	self_employed	income_annum
0	1	2	1	0	9600000
1	2	0	0	1	4100000
2	3	3	1	0	9100000
3	4	3	1	0	8200000
4	5	5	0	1	9800000

	loan_amount	loan_term	cibil_score	residential_assets_value
0	29900000	12	778	2400000
1	12200000	8	417	2700000
2	29700000	20	506	7100000
3	30700000	8	467	18200000
4	24200000	20	382	12400000

	commercial_assets_value	luxury_assets_value	bank_asset_value
loan_status			
0	17600000	22700000	8000000
1			
1	2200000	8800000	3300000
0			
2	4500000	33300000	12800000
0			
3	3300000	23300000	7900000
0			
4	8200000	29400000	5000000
0			

Exploratory Data Analysis (EDA)

- Visualize relationships between features.
- Identify trends and patterns.

1.Univariate Analysis

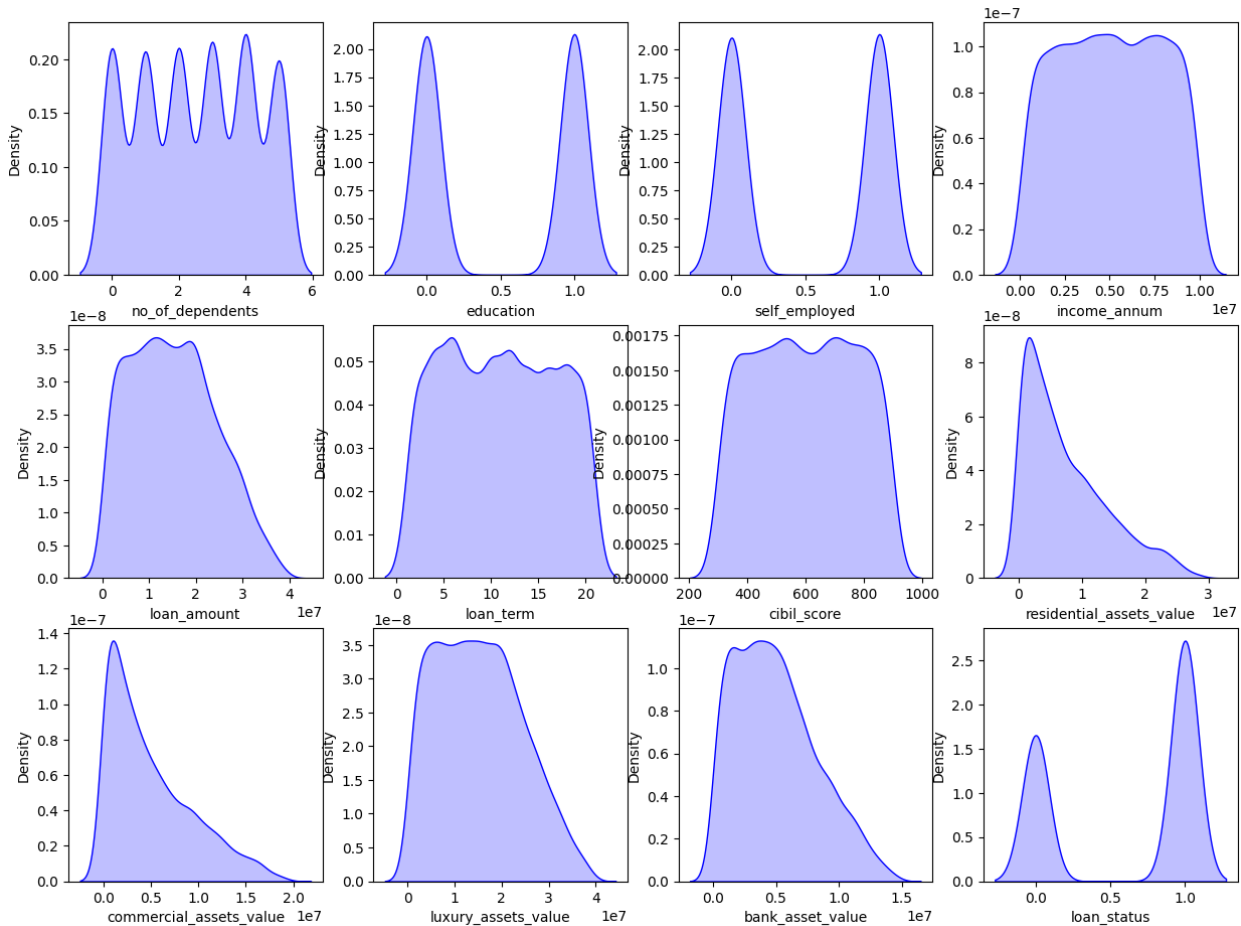
It focuses on examining a single variable at a time. It helps in understanding the distribution of the data.

```
# Univariate analysis
plt.figure(figsize=(15, 15))

numerical_columns = ['no_of_dependents', 'education', 'self_employed',
                     'income_annum', 'loan_amount', 'loan_term', 'cibil_score',
                     'residential_assets_value', 'commercial_assets_value',
                     'luxury_assets_value', 'bank_asset_value', 'loan_status']

for i in range(len(numerical_columns)):
    plt.subplot(4, 4, i + 1)
    sns.kdeplot(x=df[numerical_columns[i]], fill=True, color='b') #
    Replaced shade with fill
    plt.xlabel(numerical_columns[i])

plt.show()
```



```
print(df.columns) # List all column names
```

```
Index(['loan_id', 'no_of_dependents', 'education', 'self_employed',  
      'income_annum', 'loan_amount', 'loan_term', 'cibil_score',  
      'residential_assets_value', 'commercial_assets_value',  
      'luxury_assets_value', 'bank_asset_value', 'loan_status'],  
      dtype='object')
```

2.Bivariate Analysis

It helps in understanding the relationship between two variables, and allows us to know how they are related with each other negatively or positively.

#Bivariate Analysis

```
numerical_columns = ['no_of_dependents', 'income_annum',  
                    'loan_amount', 'loan_term',  
                    'cibil_score', 'residential_assets_value',  
                    'commercial_assets_value',  
                    'luxury_assets_value', 'bank_asset_value']  
plt.figure(figsize=(15, 12))
```

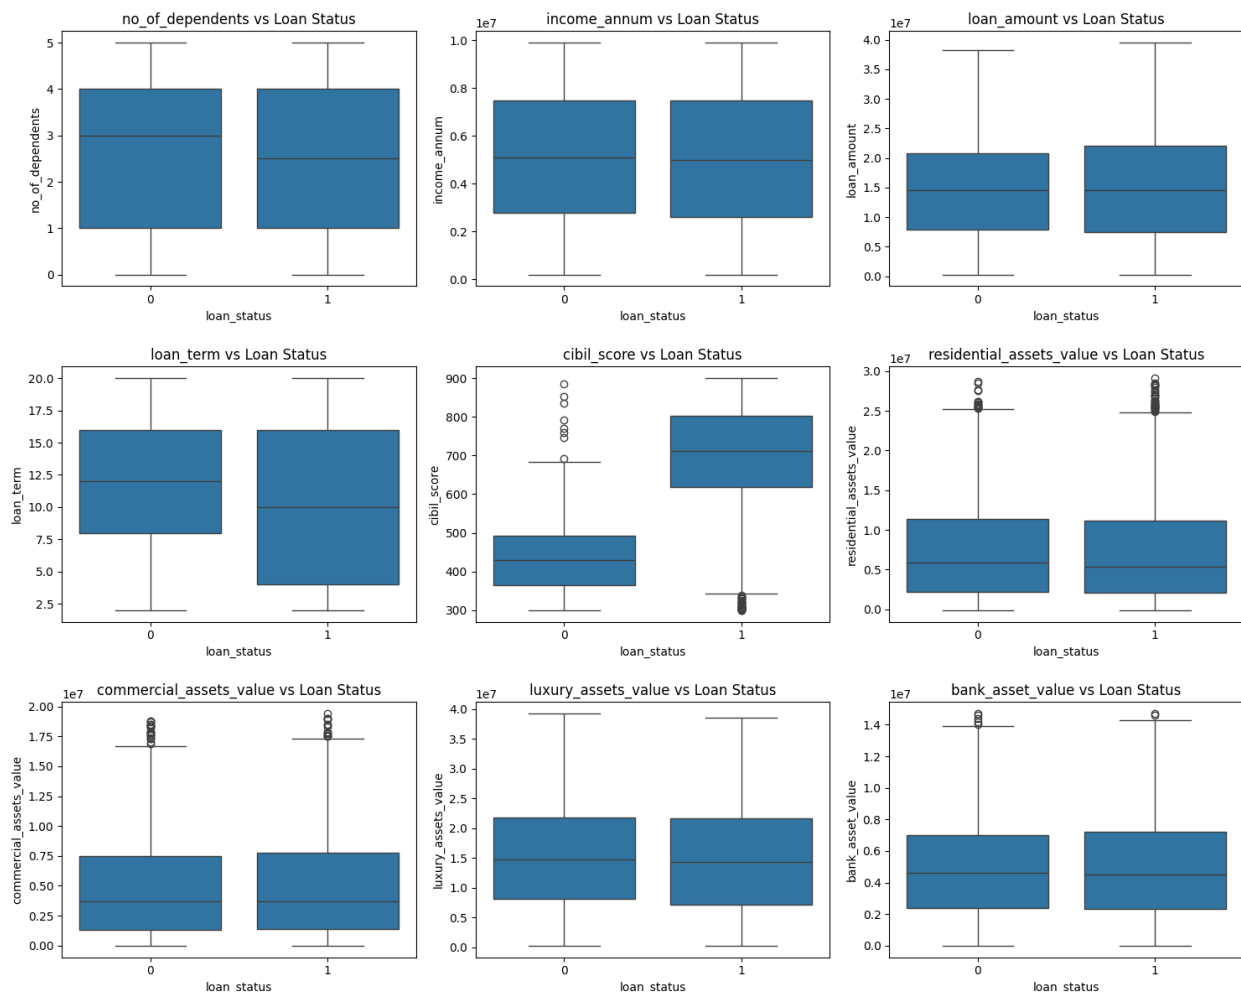


```

for i, col in enumerate(numerical_columns, 1):
    plt.subplot(3, 3, i)
    sns.boxplot(data=df, x='loan_status', y=col)
    plt.title(f"{col} vs Loan Status")

plt.tight_layout()
plt.show()

```



3.Multivariate Analysis (Relation between multiple variables)

#MultiVariate Analysis


Plot the heatmap of the correlation between each of the numerical features

```

plt.figure(figsize=(10, 8)) # Increase figure size
sns.heatmap(df[numerical_columns].corr(), annot=True, cmap='Blues',
linewidths=0.5)

plt.show()

```



Model Selection & Training:

#Adding all asset cols into one col.

```
df['total_assets'] = df['residential_assets_value'] +
df['commercial_assets_value'] + df['luxury_assets_value'] +
df['bank_asset_value']
```

#Dropping non required columns

```
df.drop(columns=['residential_assets_value',
'commercial_assets_value', 'luxury_assets_value', 'bank_asset_value'],
inplace=True)
```

```
df.head()
```

	loan_id	no_of_dependents	education	self_employed	
income_annum	\				
0	1	2	1	0	9600000
1	2	0	0	1	4100000
2	3	3	1	0	9100000
3	4	3	1	0	8200000
4	5	5	0	1	9800000

	loan_amount	loan_term	cibil_score	loan_status	total_assets
0	29900000	12	778	1	50700000
1	12200000	8	417	0	17000000
2	29700000	20	506	0	57700000
3	30700000	8	467	0	52700000
4	24200000	20	382	0	55000000

```
# Dropping the unwanted columns(Less Required Cols)
df.drop(columns=['loan_id', 'no_of_dependents', 'education',
'self_employed'], inplace=True)

#Removing Target Variable and creating the feature matrix `x`, which
will be used for training the model.
x = df.drop(columns=['loan_status'])

# Creating the Target matrix `y`, which contains target variable
y = df['loan_status']
```

Features (X) and Target (y) in Supervised Learning

- **Features (X):** Input variables used to predict the target. Examples:
 - CIBIL Score
 - Loan Amount
 - Total Assets
 - Loan Term
 - Income Annum
- **Target (y):** The output variable we want to predict, i.e., **Loan Approval Status (Approved/Not Approved)**.

Why Split the Dataset?

- **X (Features):** Helps the model learn patterns influencing loan approval.

- **y (Target):** The value the model aims to predict.
- Separating features and target allows the model to understand relationships and make accurate predictions.

Split data into training and testing sets.

```
# Divide the dataset into training (80%) and testing (20%) subsets
x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.2, random_state=42)

scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

Train the Regression Model

```
# Instantiate the Logistic Regression model
model = LogisticRegression()

# Fit the model using the training dataset
model.fit(x_train, y_train)

# Generate predictions on the test dataset
y_pred = model.predict(x_test)

# Display the first 10 predicted values
print('Predictions',y_pred[:10])

Predictions [0 1 0 1 1 1 1 0 1 0]

#Getting Model Accuracy
print('Accuracy',accuracy_score(y_test,y_pred))

print('Accuracy',classification_report(y_test,y_pred))

Accuracy 0.9063231850117096
Accuracy          precision    recall  f1-score   support

         0          0.88      0.87      0.87         318
         1          0.92      0.93      0.93         536

 accuracy          0.91         0.91         0.91         854
  macro avg          0.90          0.90          0.90         854
 weighted avg          0.91          0.91          0.91         854
```


Model Evaluation

Evaluating Model Using

- **Confusion Matrix:** Displays the true positive, false positive, true negative, and false negative values to assess classification performance.
- **ROC Curve & AUC (Area Under Curve):** Shows the trade-off between true positive rate and false positive rate, with AUC values closer to 1 indicating better performance.
- **Feature Importance (Logistic Regression Coefficients):** Highlights the most influential features in model predictions, helping in feature selection and interpretability.

#Confusion Matrix


```
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=[0,1],
yticklabels=[0,1])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



```
# ROC Curve & AUC (Area Under Curve)

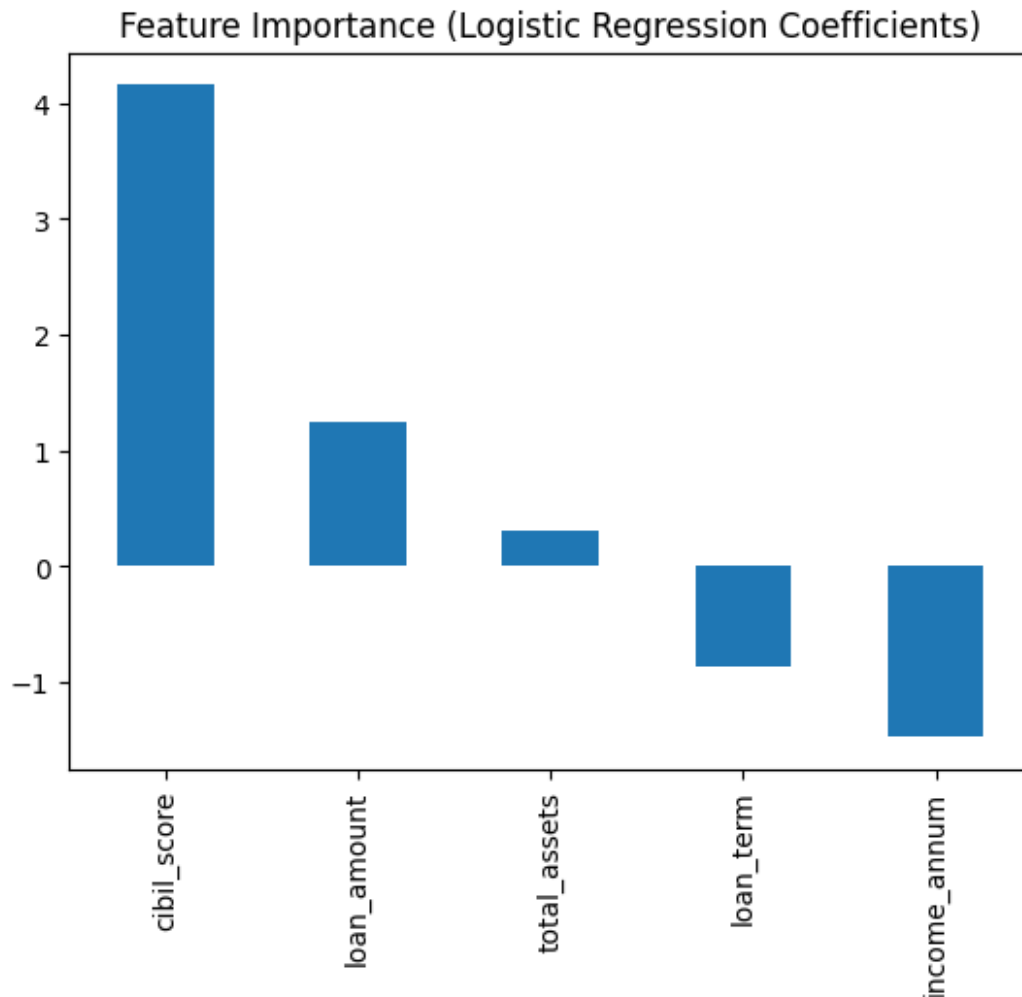
y_pred_prob = model.predict_proba(x_test)[:,-1] # Probabilities for
class 1
fpr, tpr, _ = roc_curve(y_test, y_pred_prob)
auc_score = roc_auc_score(y_test, y_pred_prob)

plt.plot(fpr, tpr, label=f"AUC = {auc_score:.2f}")
plt.plot([0,1], [0,1], 'r--') # Random guess line
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend()
plt.show()
```



```
#Feature Importance (Logistic Regression Coefficients)

importance = pd.Series(model.coef_[0], index=x.columns) # Use
original x DataFrame
importance.sort_values(ascending=False).plot(kind='bar')
plt.title("Feature Importance (Logistic Regression Coefficients)")
plt.show()
```



Model Evaluation Result

Confusion Matrix

- **True Positives (498)** and **True Negatives (276)** indicate good classification.
- **False Positives (42)** and **False Negatives (38)** should be analyzed for potential improvements.

ROC Curve (AUC = 0.97)

- A high AUC value suggests excellent model performance and strong discriminatory power.

Feature Importance

- **CIBIL Score** has the highest impact on predictions.
- **Loan Amount** also plays a significant role.
- **Income Annum** negatively influences predictions, indicating an inverse relationship.